Commentary2Highlight

Package comparison and current project report | 19 October 2026

1. What is different between the ZIP files?

The code ZIP is the lightweight source-code handoff. The handoff bundle is a reproducibility package with the code plus material to understand, rerun, and review the current milestone.

Package	What it contains	Best for
Code ZIP	16 Python scripts under src/, README.md, requirements.txt. No model weights, training labels, generated predictions, video, or virtual environment.	Reading or sharing implementation; install dependencies separately.
Handoff bundle ZIP	All code, setup notes, handoff and audit reports, selected training/review CSV and JSON files, trained PyTorch checkpoints and metrics, and selected demo outputs. No video files, Whisper/NLLB cache, or .venv.	Giving teammates the full project state and material to review or reproduce current work.

Note about the earlier handoff ZIP

I could not find an older handoff ZIP in the project workspace. The project does contain a teammate handoff report. To make the difference concrete and provide both files now, I recreated the code-only ZIP and created a clearly named handoff bundle ZIP. The bundle contents are reflected here.

2. Current pipeline and implementation status

Stage	Status	What it does now
Audio and timestamped ASR	Working, with input limits	FFmpeg extracts audio; Faster-Whisper small transcribes and detects language with estimated timestamps. Whisper translate can produce English.
Text preparation and NLP cues	Working baseline	Project normalization and event phrases retain negation, hypothetical, retrospective, live-action, and emphasis cues.
Candidate events and frames	Partial	Rules propose goal, penalty/VAR, card, shot/chance, and set pieces. Lexical frames are not full semantic-role labeling.
Project-trained DL ranker	Implemented; experimental	Project-fitted TF-IDF feeds a PyTorch MLP trained from random initialization. Multi-task MLP and BiGRU with temporal attention are implemented but experimental.
Entities/coreference	Partial	Exact aliases and cautious short-window pronoun links work if a verified roster is supplied. No trained player-name/nickname model.
Peak grouping and exports	Working baseline	Nearby candidates are grouped and exported as timestamped Markdown, CSV, and JSON. No edited video clips yet.
Evaluation	Needs independent labels	Whole-match-held-out exploratory evaluation spans seven matches, but mixes weak SoccerNet alignment and limited human review.

3. Models and what they contribute

Pretrained input components

Faster-Whisper small supplies multilingual speech recognition, language detection, and timestamps; English translation may use Whisper translate. NLLB-200 distilled 600M is an optional text-translation route. These are pretrained input tools.

Team-built NLP

Normalization, football event phrase rules, transparent discourse cues, TF-IDF feature fitting, and event-frame templates are project code.

Team-trained DL

A PyTorch MLP learns event-related vs routine scores from project-built TF-IDF. A multi-task MLP adds broad event-family predictions. A BiGRU with temporal attention processes ordered transcript windows. These downstream models are trained by the team; they do not fine-tune Whisper or NLLB.

Current reported evaluation

Project notes report seven-match whole-match-held-out mean F1 of 0.655 for the TF-IDF MLP, 0.553 for the TF-IDF linear baseline, and 0.597 eventness macro-F1 for the experimental sequence model. Mixed weak labels and small reviewed data mean these are exploratory, not broad independent evidence.

4. Demonstrated outputs and known boundaries

• **Timestamped output:** Crystal Palace vs Manchester City three-minute sequence demo produced 81 transcript segments and two candidates near 00:01:23 and 00:02:50. Structured results are included in the handoff bundle.

• **Language checks:** Spanish, French, and Russian examples exist. Translation and player-name spelling vary; the system is not validated for every language or broadcast.

• **Video/audio:** inventory lists five SoccerNet MKVs; three have usable audio and two Manchester City-Newcastle files decode to no audio samples. Check each input before transcription.

• **Not reliable/complete:** full coreference, full SRL, learned player identity/nicknames, visual event detection, automatic clip cutting, and broad cross-language benchmark evaluation.

5. Best next work

1. Finalize human labels for candidate moments and routine controls; check duplicates before merging into training.

2. Keep a complete match untouched for final evaluation and add independent matches and languages.

3. Improve event extraction with reviewed positives and negatives; separate event-related from subjective-highlight labels.

4. Build verified team/player aliases for selected fixtures; measure name corrections instead of claiming general recognition.

5. Compare translation routes against source audio and check timestamp offsets. Defer clip cutting until timestamps are reviewed.

6. Keep the TF-IDF MLP as the primary project-trained ranker; report multi-task and BiGRU-attention as experiments until they improve held-out results.

6. Suggested team explanation

“We have a working prototype that extracts audio, creates timestamped multilingual transcripts, identifies event-related commentary, and ranks candidate moments with our own NLP features and PyTorch models trained from project data. Whisper and optional translation are pretrained input tools. Our event rules, TF-IDF representation, MLP, and sequence model are ours. Remaining work is better reviewed labels, reliable cross-match evaluation, name handling, and timestamp validation. The current deliverable is timestamped highlight commentary; automatic video cutting is future work.”

Prepared from the project README, teammate handoff, progress review, pipeline/data audit, and model metrics as of 19 October 2026.